

Traffic Management Applications for Stateful SDN Data Plane

Carmelo Cascone[^], Luca Pollini*, Davide Sanvito*, Antonio Capone[^]

ANTLab - Politecnico di Milano[^]
Dipartimento di Elettronica, Informazione
e Bioingegneria



CNIT*
Consorzio Nazionale Interuniversitario
per le Telecomunicazioni




Supported by EU project:



Goal

- **Highlight shortcomings of current SDN-OpenFlow paradigm**
- **Present a new “stateful” data plane model**
- **Motivate this need with 2 application examples**
 - Failure recovery
 - Forwarding consistency


OpenFlow recap



Centralized control: we know the pros but...

- **Control latency**
 - Switch-controller RTT
 - Controller processing
- **Signaling overhead**
 - First packet to the controller (Internet dominated by very short flows)
 - Flow statistics gathering


Example: failure recovery in OpenFlow (1)



“Fast-failover”:
Local reroute based
on port status
(OpenFlow 1.1+)


Weak! What if a
local reroute is not
available?

Example: failure recovery in OpenFlow (2)



- Can rely on controller intervention, but:
 - Long recovery latency (> 50ms)
 - detection + signaling + flow table update
 - Failure of control channel
 - Signaling congestion (e.g. multiple failures, disasters)

Towards a new behavioral data plane model




Easier said than done

- We need a switch abstraction and API which is...
 - High performance: control tasks executed at wire-speed (packet-based events)
 - Platform-independent: consistent with vendors' needs for closed platforms
 - Low cost and immediately viable: based on commodity HW
- Apparently, far beyond OpenFlow switches...
- Our finding: much closer to OpenFlow than expected

Our approach: OpenState


- Idea: forward packets based on “flow states”
 - Maintained by the switch
 - Autonomously updated as a consequence of local events (i.e. match, timers)

- FSM-like forwarding model
- Minimal extension to OpenFlow




- [CCR '14] G. Bianchi, M. Bonola, A. Capone, C. Cascone, “OpenState: programming platform-independent stateful OpenFlow applications inside the switch”, ACM SIGCOMM Comp. Comm. Rev., April 2014
- [HPSR '15] S. Pontarelli, M. Bonola, G. Bianchi, A. Capone, C. Cascone, “Stateful OpenFlow: Hardware Proof of Concept”, IEEE High Performance Switching and Routing, July 2015

OpenState: 2 table approach




Flow key extractors

- Used to match/access the state table
 - Lookup or update phase
- Scope = ordered list of header fields
 - E.g. {ip_src} → 32 bit flow key
 - E.g. {eth_src, eth_dst} → 96 bit flow key




State table

- **Exact match on flow key**
 - Efficient implementation in RAM (vs. TCAM)
- **DEFAULT state if table miss**
- **Optional timeouts**
 - Idle or hard: equivalent to OpenFlow
 - $\leq 1\text{ms}$ granularity
 - Rollback state when timeout expires
 - Configured by `set_state()` action



Pipeline configuration

Tables are **stateless at switch boot**.
The controller can then configure one
or more tables as stateful.




Open source: <http://www.openstate-sdn.org>

- **Running code: softswitch + controller**
 - Based on CPqD ofsoftswitch13, RYU
 - Initial support to Open vSwitch based on “learn()” action
- **Protocol specification**
 - OpenFlow 1.3 Experimenter Extension (PDF available)
- **Mininet-based application examples**
 - MAC learning, port knocking firewall, failure Recovery, DDoS detection and mitigation, load balancing
- **Download & try!**

Failure recovery

Failure recovery with OpenState


- Tags (e.g. MPLS labels) used to distinguish between **different forwarding behaviors**
- Upon failure, packets are “**bounced back**” with special tag
 - until matched against a node able to respond to that specific failure
- **Periodic probe** to re-establish forwarding on the primary path



- ➔ **No extra signaling/packet loss after failure detection**
- ➔ **Controller not involved (besides initial provisioning)**


Behavioral model (FSM)

- Each flow (lookup-scope) has an associated state (tag)
 - **0 (default)** → all good, forward on primary path
 - **Fi** node i unreachable → forward on detour i -th
 - **Pi** node i must be probed → send 1 probe to node i



Failure recovery Example

Normal conditions (no failures)



lookup-scope=[eth_src, eth_dst]
update-scope=[eth_src, eth_dst]

L2 flows

State table


Key	State
...	...
...	...
* (any)	0

Flow table

Match	Instructions
src=1, dst=6, state=0	fwd(3)
...	...
...	...

Failure recovery Example

Packets “bounced back” in case of failure




Match	Instructions
src=1, dst=6	Group(1)
...	...
...	...

Group table


ID	Type	Action buckets
1	FAST-FAILOVER	<output(2)>, <push_tag(F4), output(1)>,
...

Failure recovery Example

State transition at a pre-determined reroute node




Match	Instructions
...	...
src=1, dst=6, state=0	fwd(3)
src=1, dst=6, tag=F4	set_state(F4, hard_to=10s, hard_rollback=P4) fwd(7)
...	...




Failure recovery Example

Detour path enabled




Match	Instructions
...	...
src=1, dst=6, state=F4	push_tag(F4), fwd(7)
...	...
...	...



Failure recovery


Example

State hard timeout to generate probe packets




Failure recovery Example

Primary path re-established




Match	
...	...
...	...
...	...
...	...
tag=P4	set_state(0), drop()



Failure recovery Example

Failure solved



Match	Instructions
src=1, dst=6, state=0	fwd(3)
...	...
...	...

Load balancing


Load balancing in OpenFlow

- **OpenFlow SELECT group entry**
 - Packets forwarded using only one of multiple defined action buckets
 - Implementation left out to vendors (e.g. round robin, hash-based, etc)
- **Usually implemented with ECMP-like hash-based schemes**
 - Can't decide on which header fields
 - Two or more elephant flows can collide on their hash, using the same path, hence creating a bottleneck
 - **Current OF solutions:**
 - reactive allocation (first packet to controller)
 - detection and relocation based on periodic flow statistic gathering

Better idea: flowlet-based load balancing

- Originally introduced in FLARE (2007)*

- Based on the idea of **switching bursts of packets** (flowlets) instead of pinning the whole flow to one path
- **No packet reordering** if the idle time between bursts is larger than the maximum delay difference between parallel paths
- No need to worry about elephant flows (burden shared among all paths)




No packet reordering if $idle_time > |delay1 - delay2|$

* S. Kandula, D. Katabi, S. Sinha, and A. Berger, “FLARE: Dynamic load balancing without packet reordering”.
ACM SIGCOMM Computer Communication Review, 2007.

OpenState-based implementation

- States used to **distinguish between consecutive bursts/instances** of the same flow
- State idle timeouts to define the **lifetime of a forwarding decision**
 - sub-RTT scales for flowlet switching



lookup_scope=[ip_src, ip_dst, tcp_src, tcp_dst]
update_scope=[ip_src, ip_dst, tcp_src, tcp_dst]

State table

Key	State	Timeouts
A,B,x,y	1	idle_to=δ
...
*	0	n/a

Flow table


Match	Instructions
ip_dst=A, state=0	group(1)
ip_dst=B, state=0	group(2)
state=1	output(1)
state=2	output(2)
...
state=N	output(N)

Group table

Group ID	Type	Action buckets
1	SELECT	<set_state(1, idle_to=δ), output(1)>, <set_state(2, idle_to=δ), output(2)>, ...
2	SELECT	...


Example results: failure recovery

- **OF:** OpenFlow-based reactive approach, controller establishes backup path (with different switch-controller RTTs)
- **OS:** OpenState-based approach, packets bounced back upon failure



Optimal routing that minimizes bounce path based on:


A. Capone, C. Cascone, A. Q. Nguyen, and B. Sansò. "Detour planning for fast and reliable failure recovery in SDN with OpenState". In IEEE Design of Reliable Communication Networks (DRCN), March 2015



Example results: load balancing

- **OF:** controller-based reactive approach, new connections allocated by controller
- **OVS:** same as OF, but with faster switch (Open vSwitch)
- **OS:** OpenState-based approach

12ms
switch-controller RTT



Conclusions

- **New stateful data plane model → OpenState**
 - Control «decided» at controller, «execution» delegated to switches' data plane)
- **Running code available at: <http://www.openstate-sdn.org>**
 - Openflow 1.3 extension
- **Failure recovery**
 - Switches pre-loaded with backup routing
 - MPLS labels use to perform failure signaling/path probing
 - Almost 0 packets lost after failure detection
- **Load balancing**
 - Can implement flowlet-based scheme
 - No need for elephant flows handling
 - Controller initially configure group table with optimal state idle timeouts



<http://www.beba-project.eu>

- Started January 2015
- Technical plans:
 - Propose OpenState for standardization
 - SW switch acceleration + HW prototype
 - Advanced security, forwarding and monitoring applications
 - Data plane verification
 - Real field large scale experimentation



Thanks!

carmelo.cascone@polimi.it