

FlowBlaze

Stateful Packet Processing in Hardware

Salvatore Pontarelli, **Roberto Bifulco**, Marco Bonola, Carmelo Cascone,
Marco Spaziani, Valerio Bruschi, Davide Sanvito, Giuseppe Siracusano,
Antonio Capone, Michio Honda, Felipe Huici, Giuseppe Bianchi



NEC



UNIVERSITA' DEGLI STUDI
DI ROMA TOR VERGATA

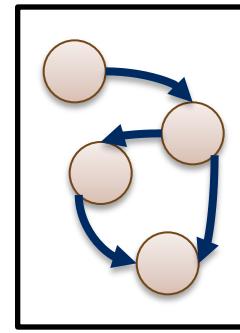
Cnict

consorzio nazionale
interuniversitario
per le telecomunicazioni



This work has received funding from the European Union's Horizon 2020 research and innovation programme under grant agreements No 761493 ("5GTANGO") and No.762057 ("5G-PICTURE")

*Now at ONF



Roberto

State
Machines

State Machines

State machines to the rescue of complex forms

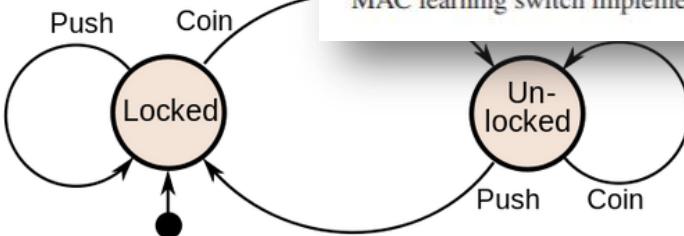


Łukasz Makuch [Follow](#)
May 6, 2017 · 6 min read

Tackling UI State Machine



Carlos Galarza [Follow](#)



tion. Finally, to show Kinetic's generality, we present a MAC learning switch implementation.

Kinetic, NSDI'15

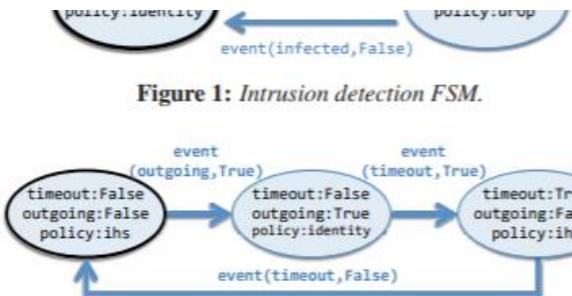


Figure 1: Intrusion detection FSM.

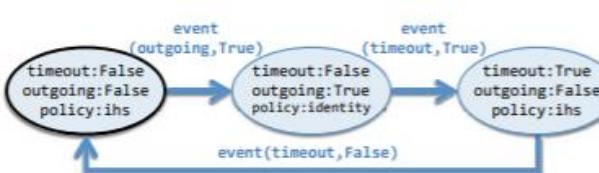


Figure 2: Stateful firewall FSM.

The Rise Of The State Machines

QUICK SUMMARY ↗ The UI development became difficult in the

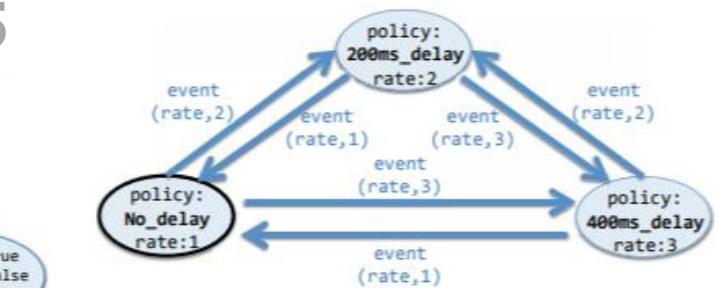
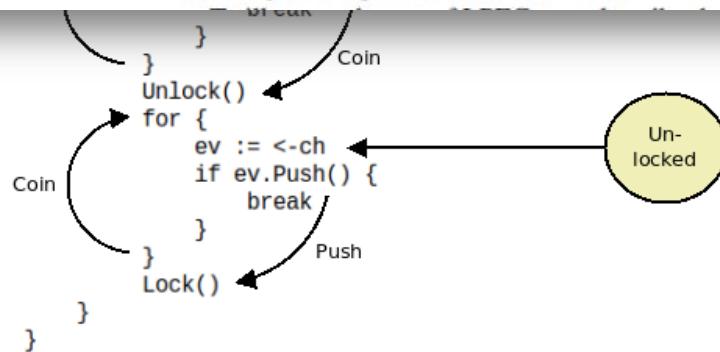


Figure 3: Data usage-based rate limiter FSM.

packets (e.g., all packets from the same host, in the case of the previous example). Each group of packets has a separate FSM instance; packets in the same group will always be in the same state. We call such a group of packets a *located packet equivalence class (LPEC)*.

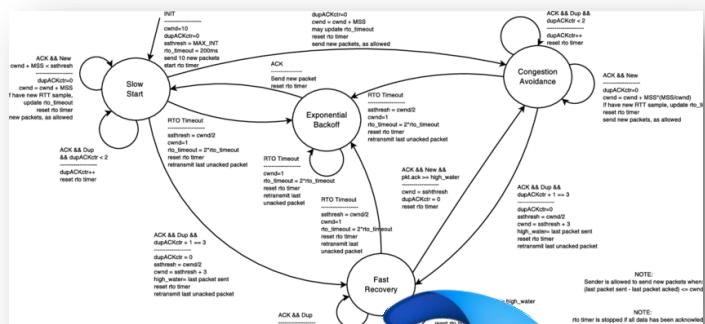


Programmable NICs

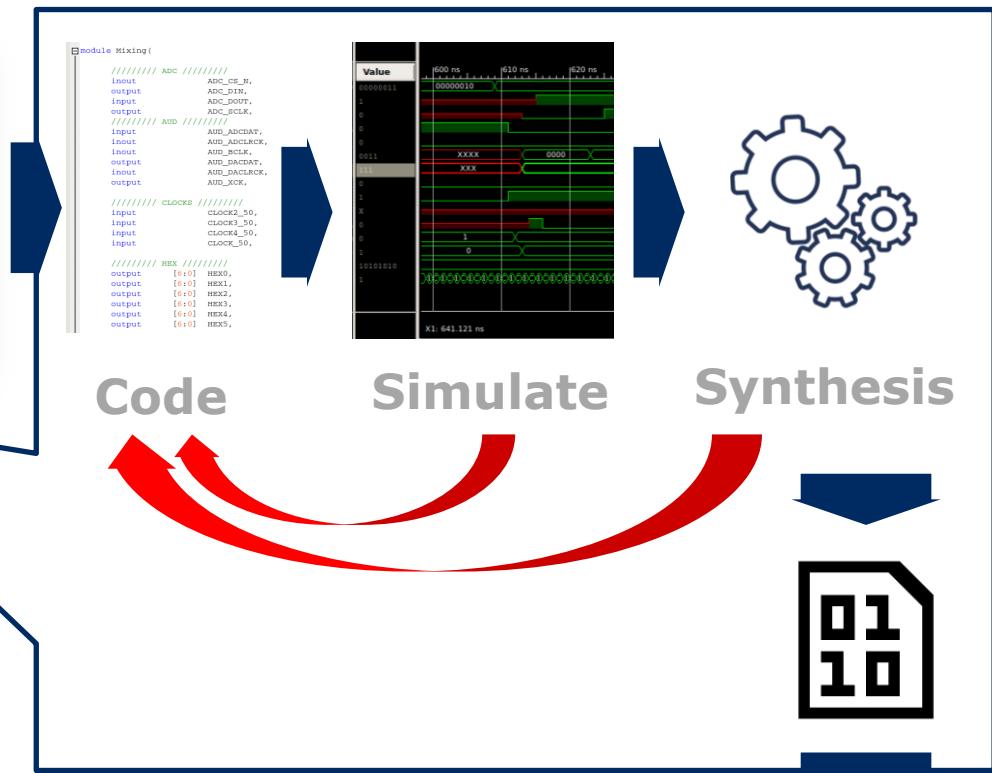
NICs with programmable ASICs, SoC, FPGAs...

E.g., Microsoft [AccelNet NSDI '17, NSDI '18]

Network Function Logic

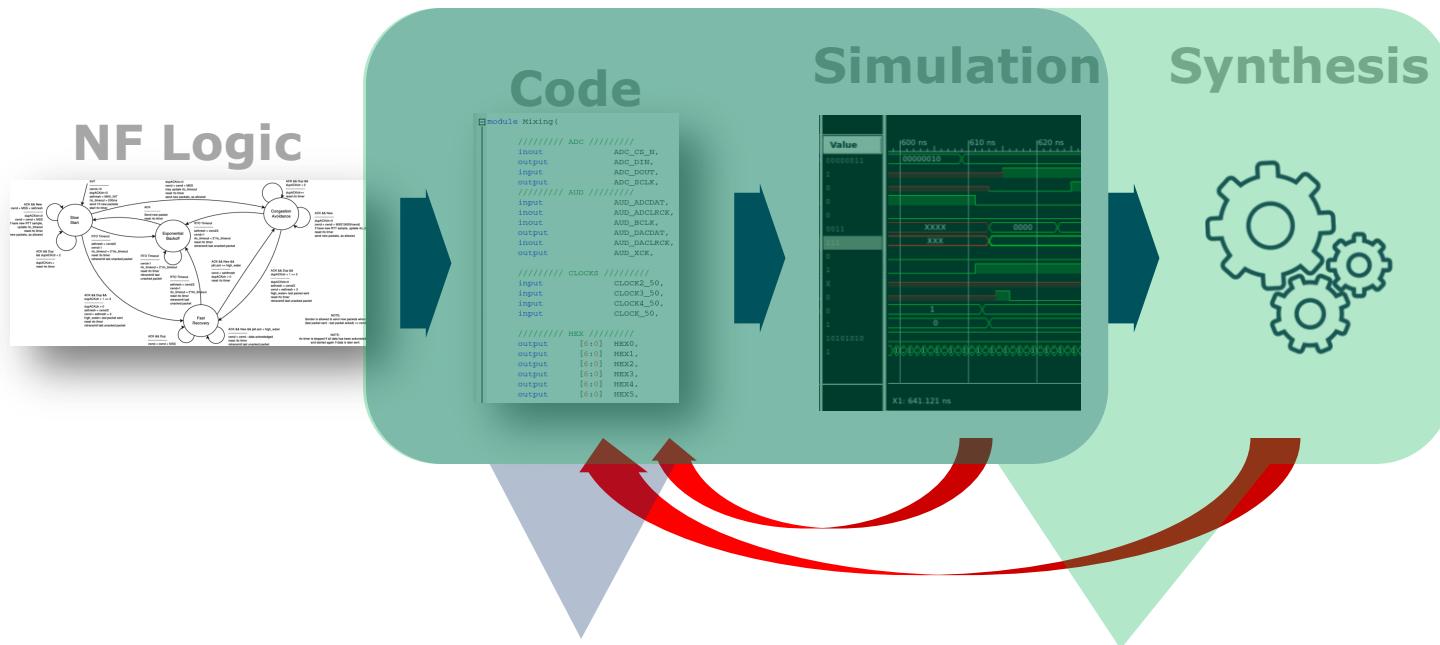


Programming them is Hard!



NetFPGA SUME

Making programming easier



High-level Synthesis

Faster programming
Expressive
Hardware expertise

ClickNP [Sigcomm '16],
Emu [ATC '17]

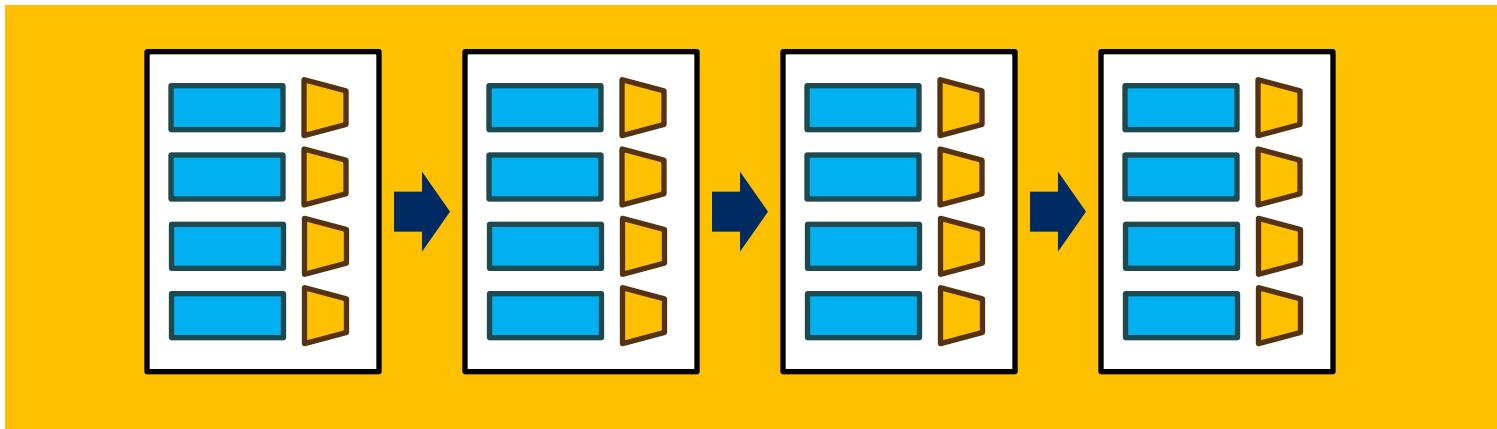
Match-Action Abstraction

Faster programming
NF Logic focused
Limited support for state

P4 [CCR '14],
Domino [Sigcomm '16]

Match-Action Abstraction Limitations

Match-Action pipeline



State in tables

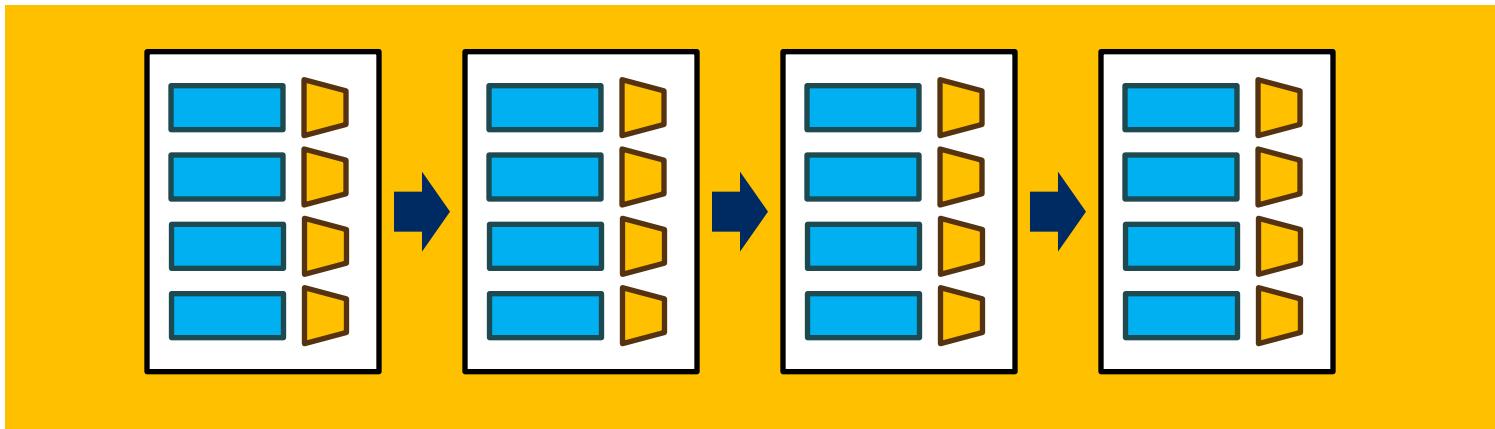
large
read only
(wr from cplane)

State in registers

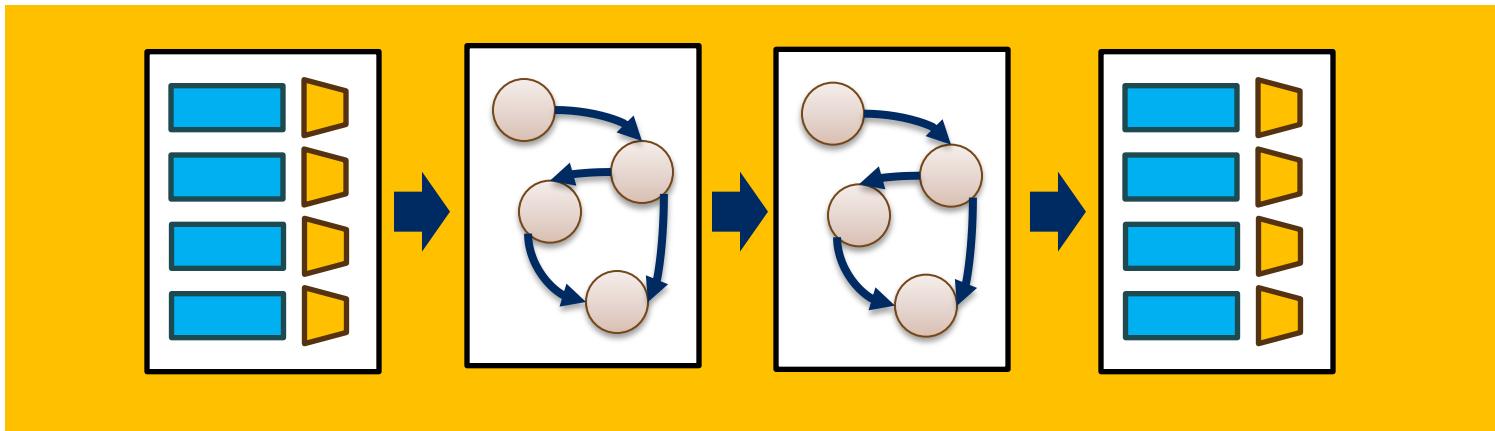
small
read/write

Extending Match-Action abstractions

Match-Action pipeline

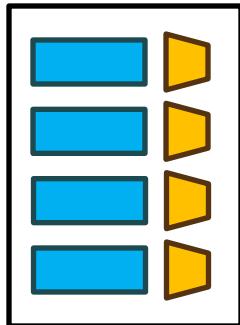


FlowBlaze



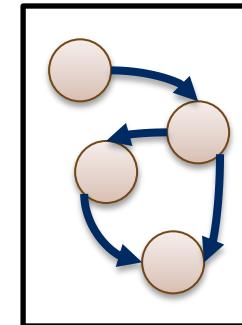
Match-Action vs Finite State Machine (FSM)

**if match
then action**

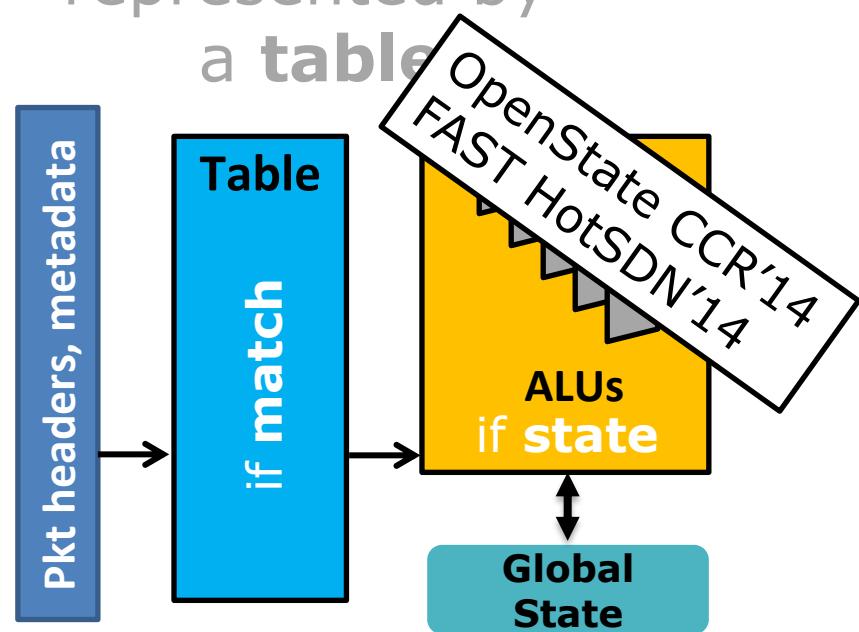
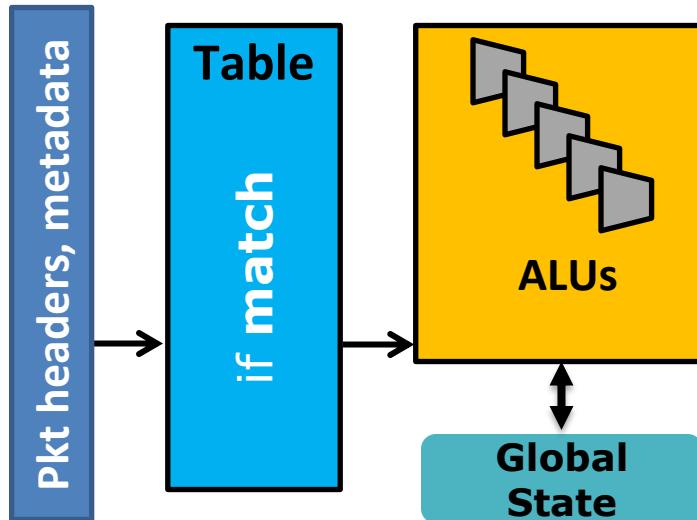


a table

**if (match, state)
then action**

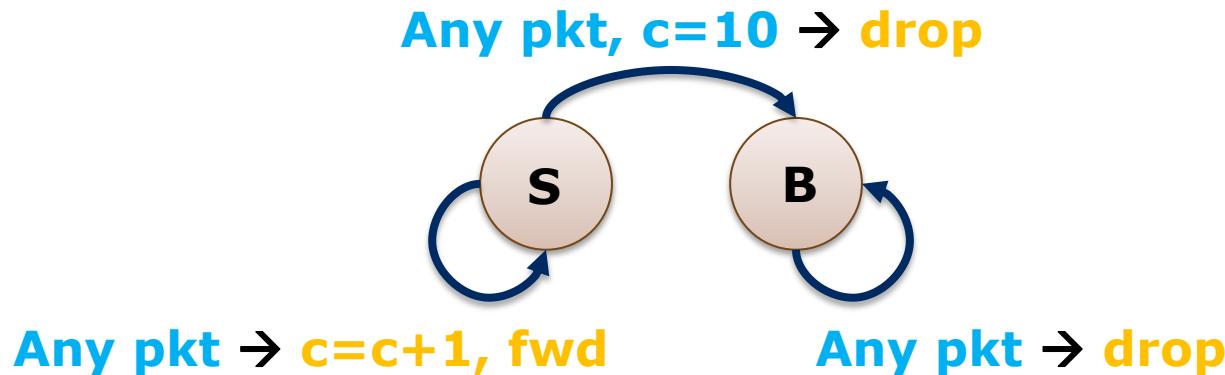


**represented by
a table**



Multiple state machines?

Example: Drop a flow after its 10th packet

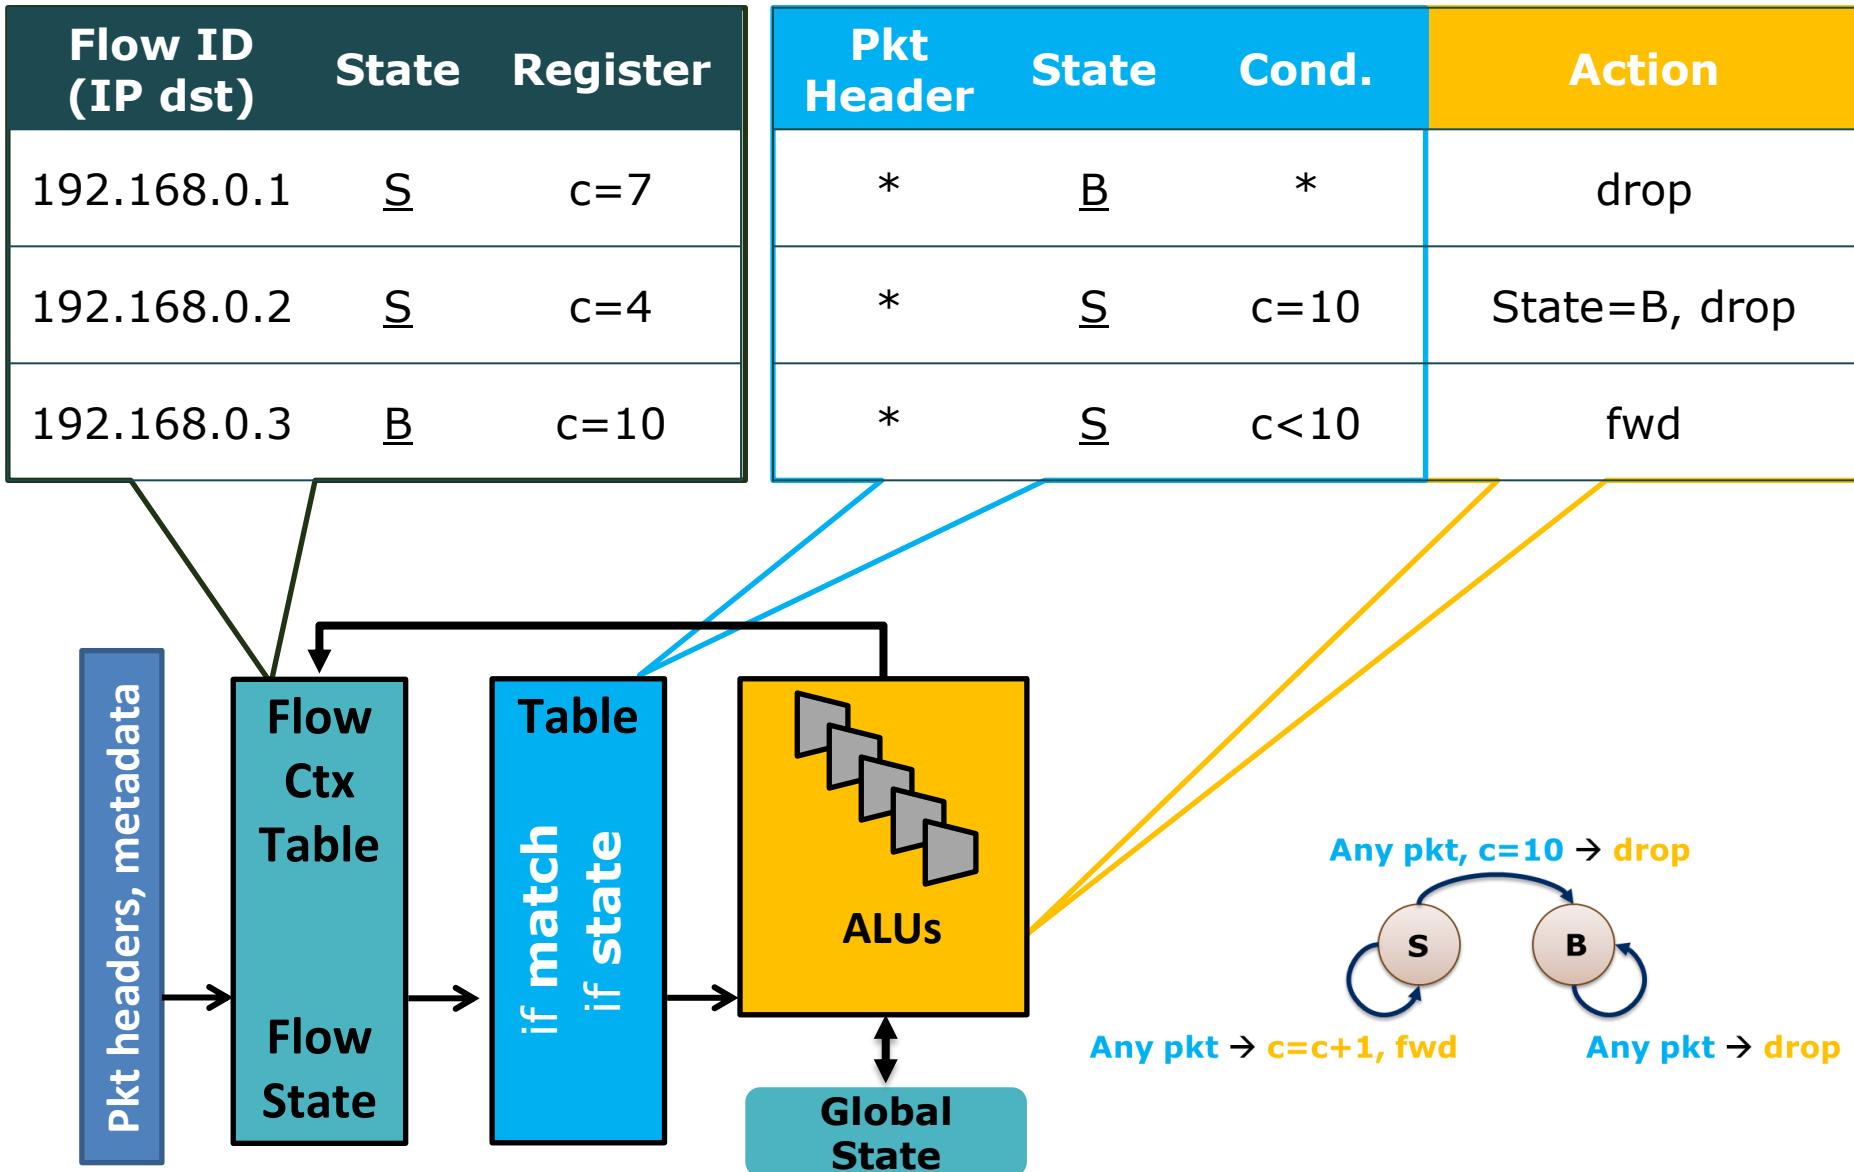


Flow ID	State
IPdst = 192.168.0.1	S
IPdst = 192.168.0.2	S
IPdst = 192.168.0.3	B

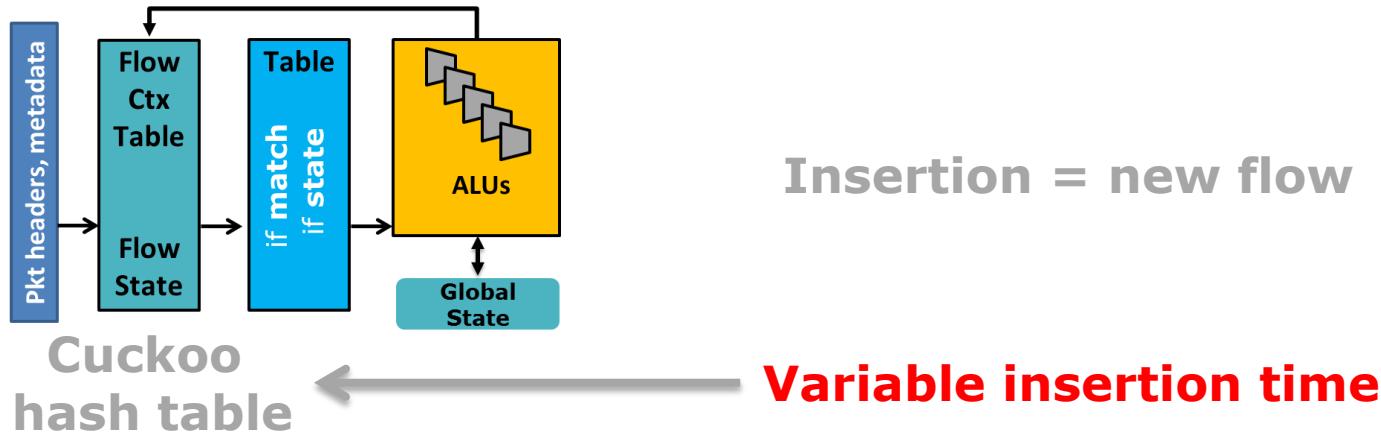
Each flow's FSM evolves on its own

Per-flow state is common in network functions

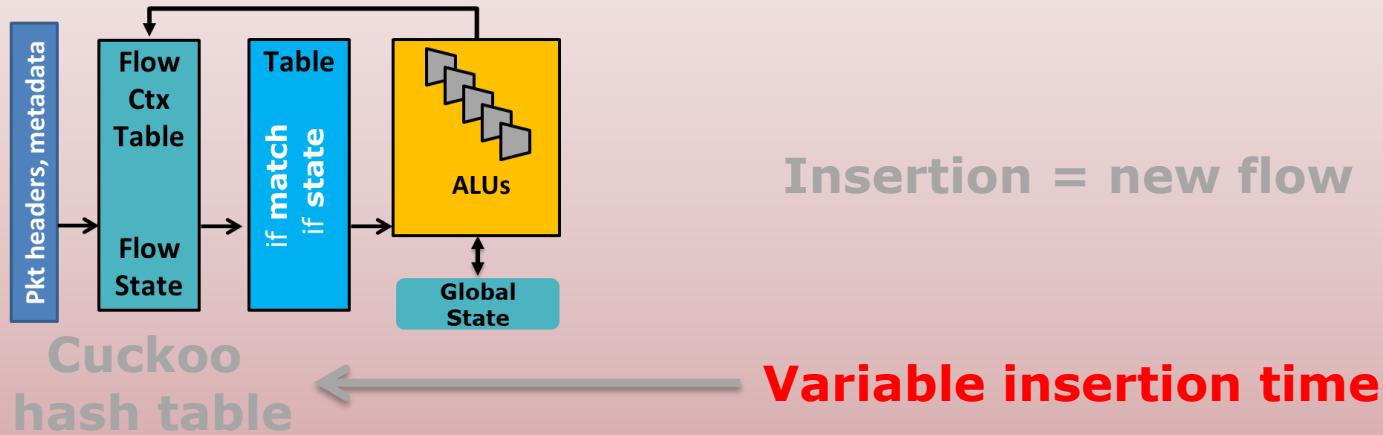
Introducing per-flow state



Insertion in the flow context table

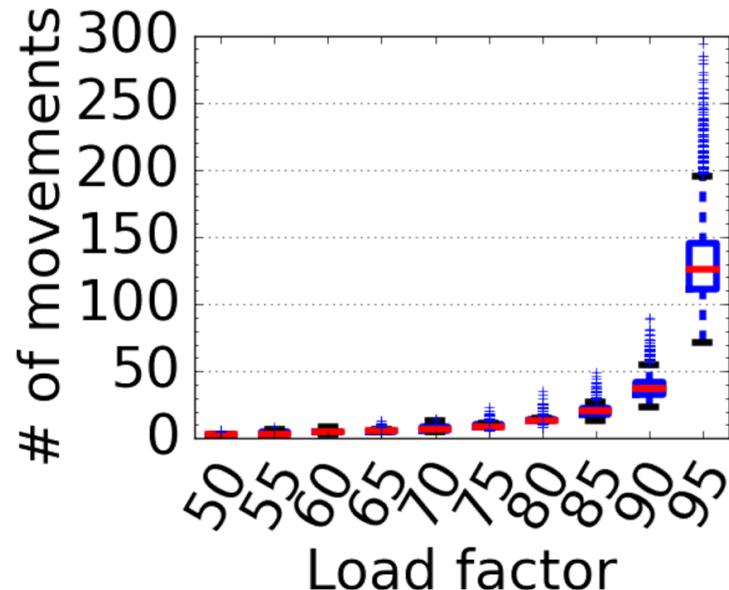


Insertion in the flow context table

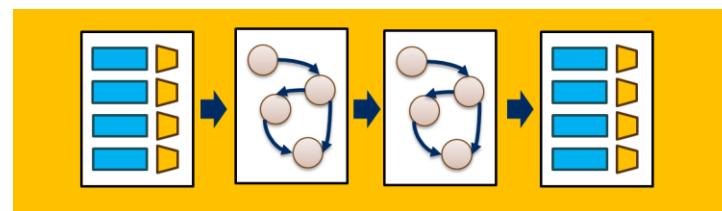


Flow table: Cuckoo hash

Efficient
Constant lookup-time
Variable insertion-time



**Waiting for
Insertion!!**



**Throughput
reduction**

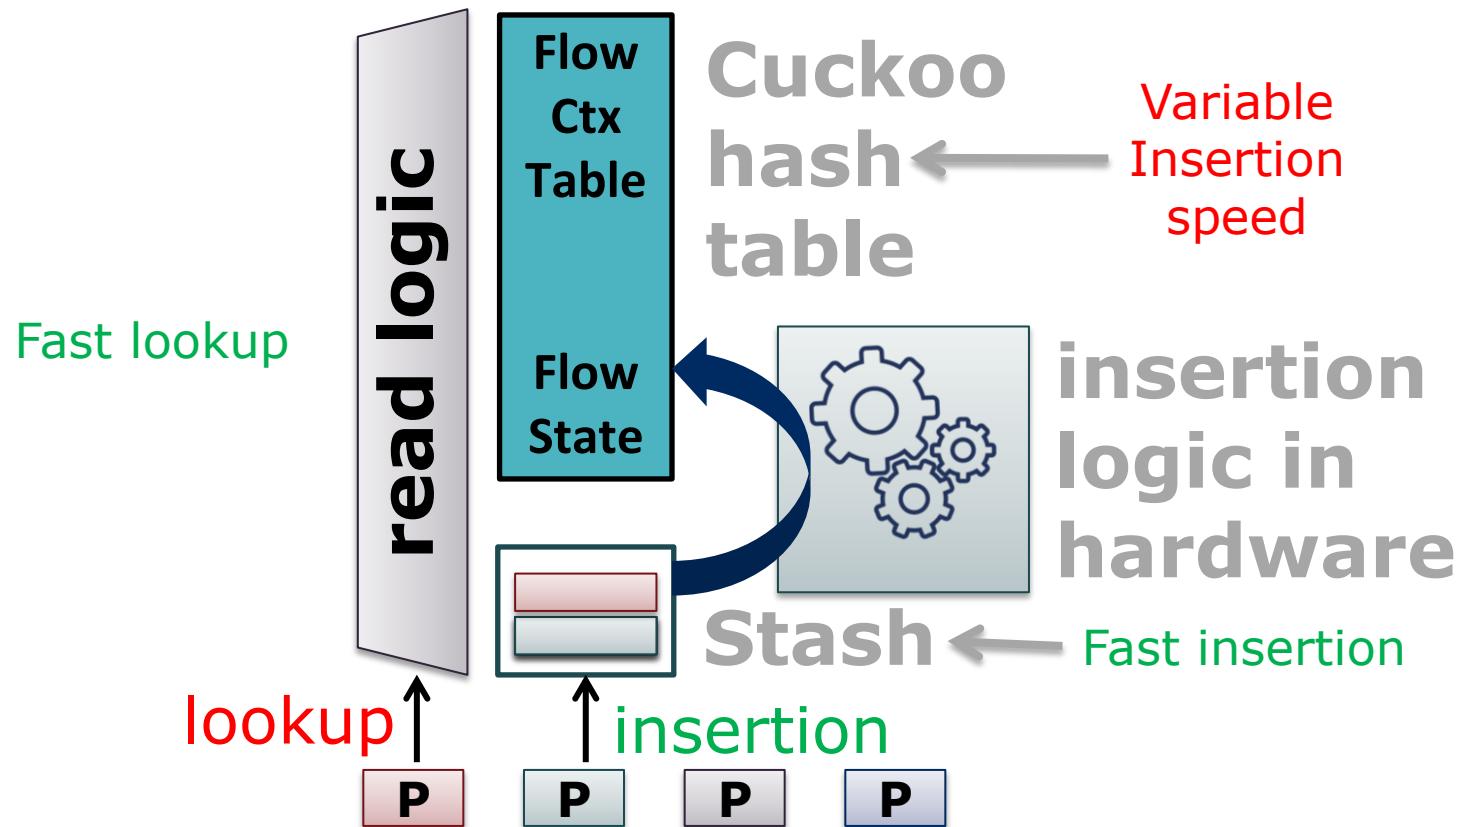


Latency increase

Flow context insertion handling

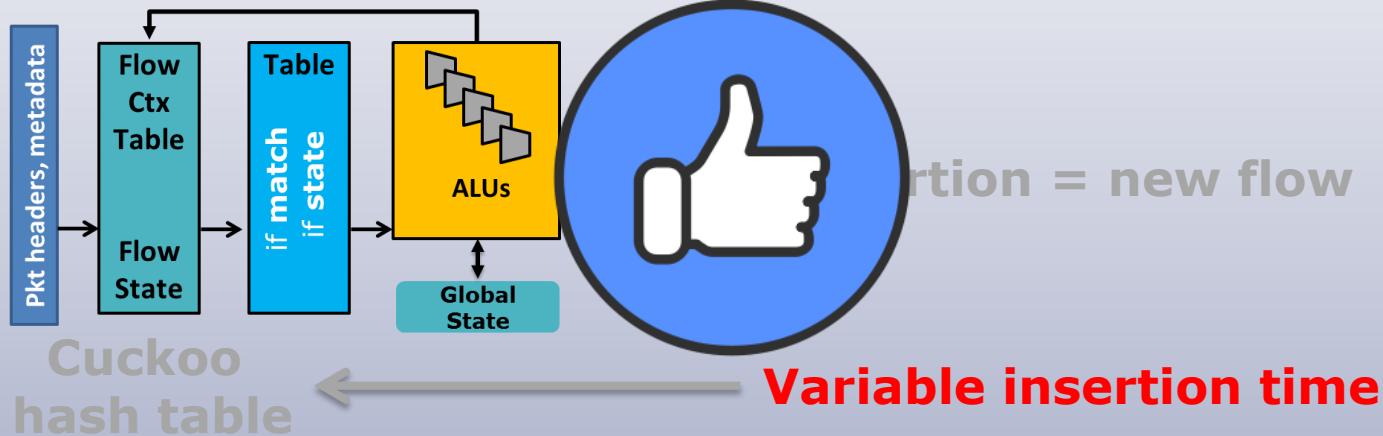
lookup time
scales with
pkt arrival rate

insertion time
scales with
flow arrival rate

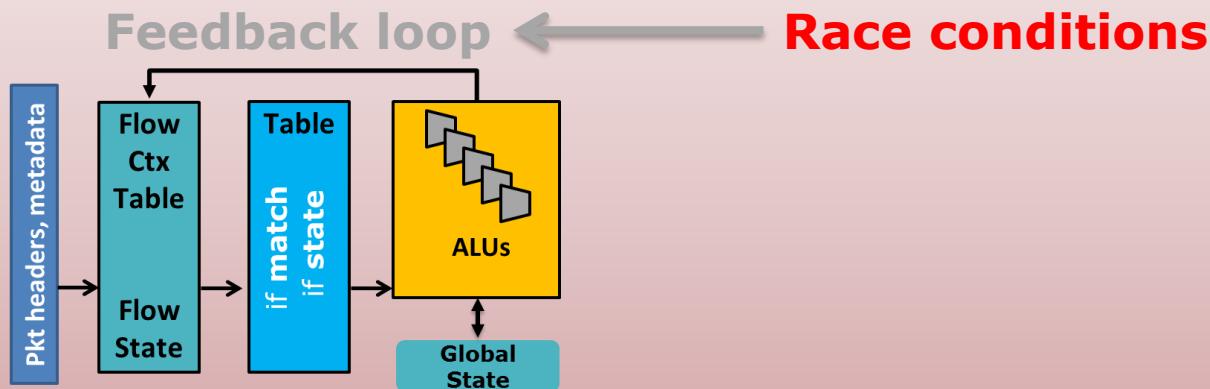


Implementation issues

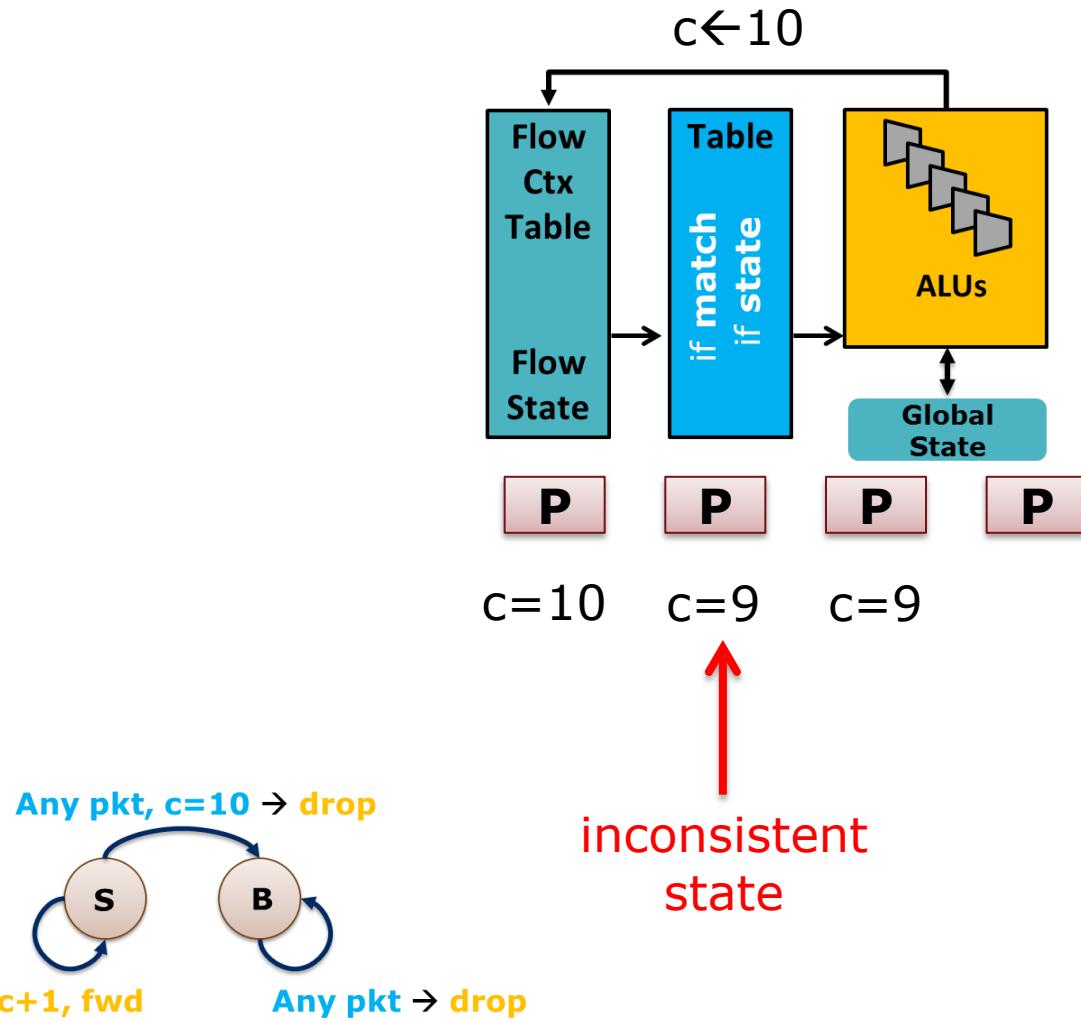
Insertion in the flow context table



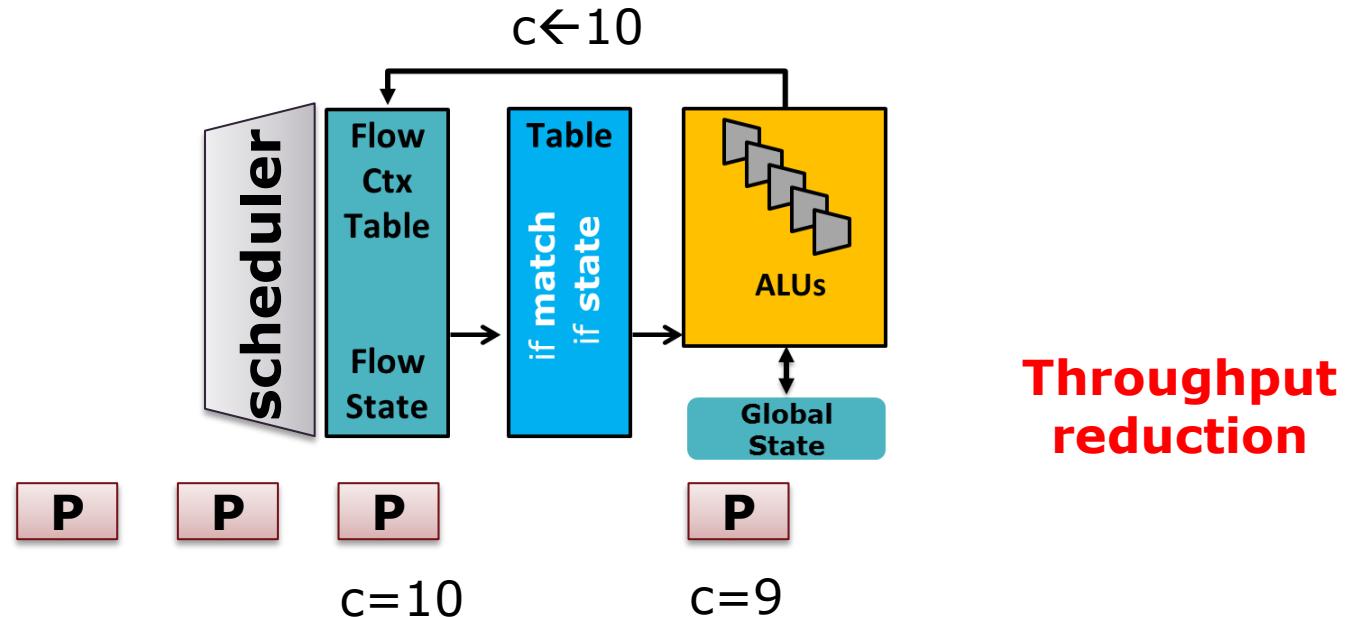
State update latency



Avoiding race conditions



Avoiding race conditions



Latency increase

Any pkt, $c=10 \rightarrow$ drop

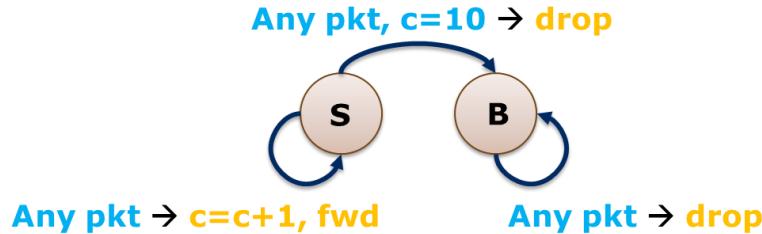
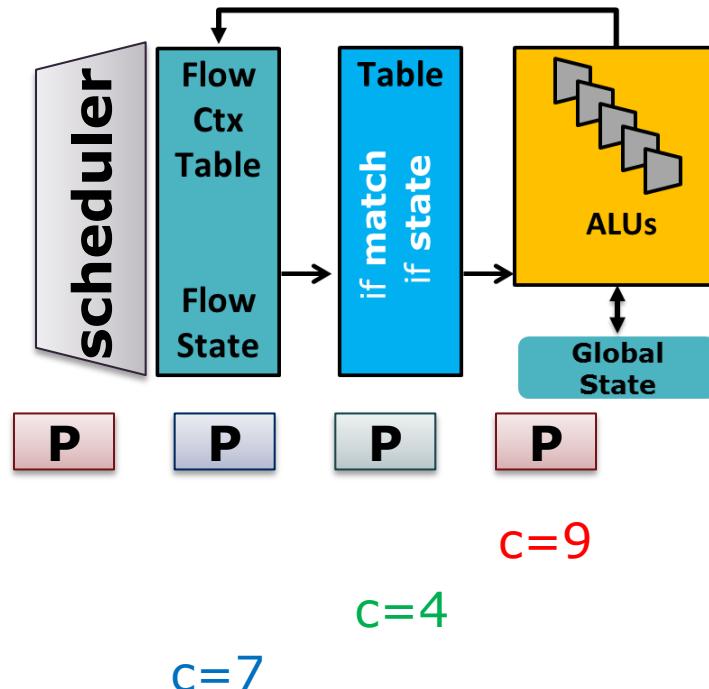


Any pkt $\rightarrow c=c+1, \text{fwd}$

Any pkt \rightarrow drop

Avoiding race conditions

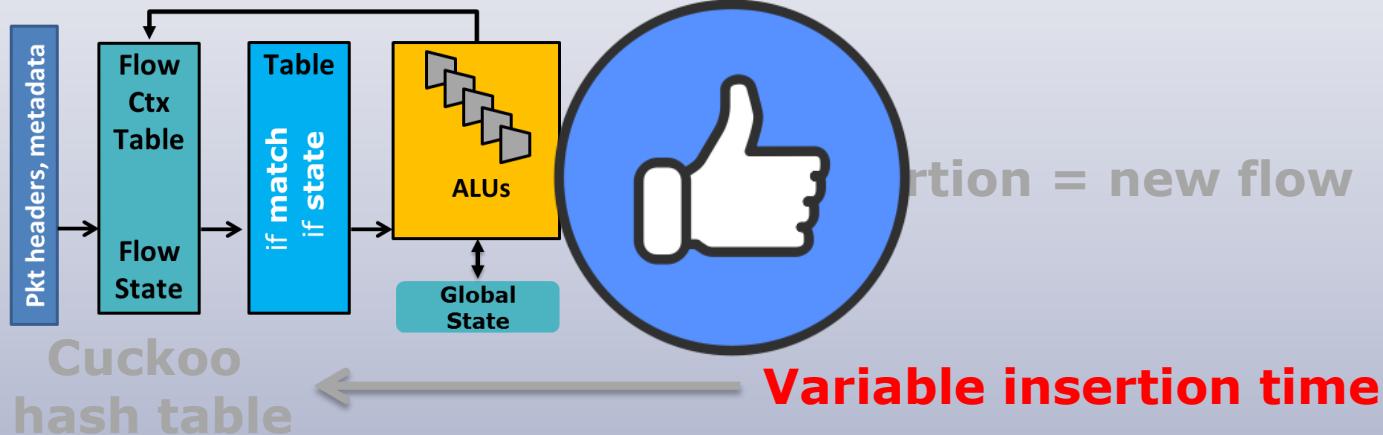
Lock pipeline
for packets
from the
same flow



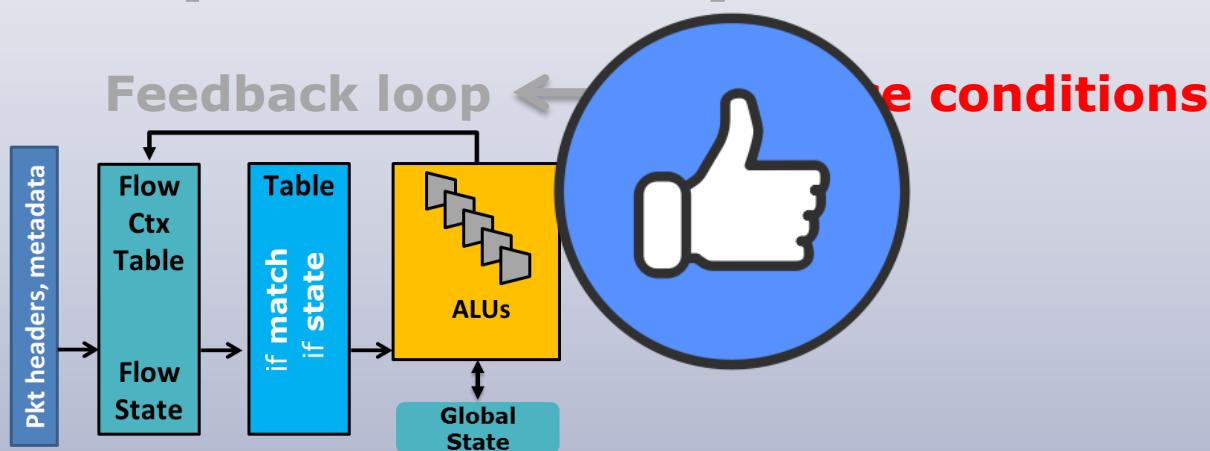
Performance
degradation only
in unlikely cases

Implementation issues

Insertion in the flow context table



State update latency



Does it work?

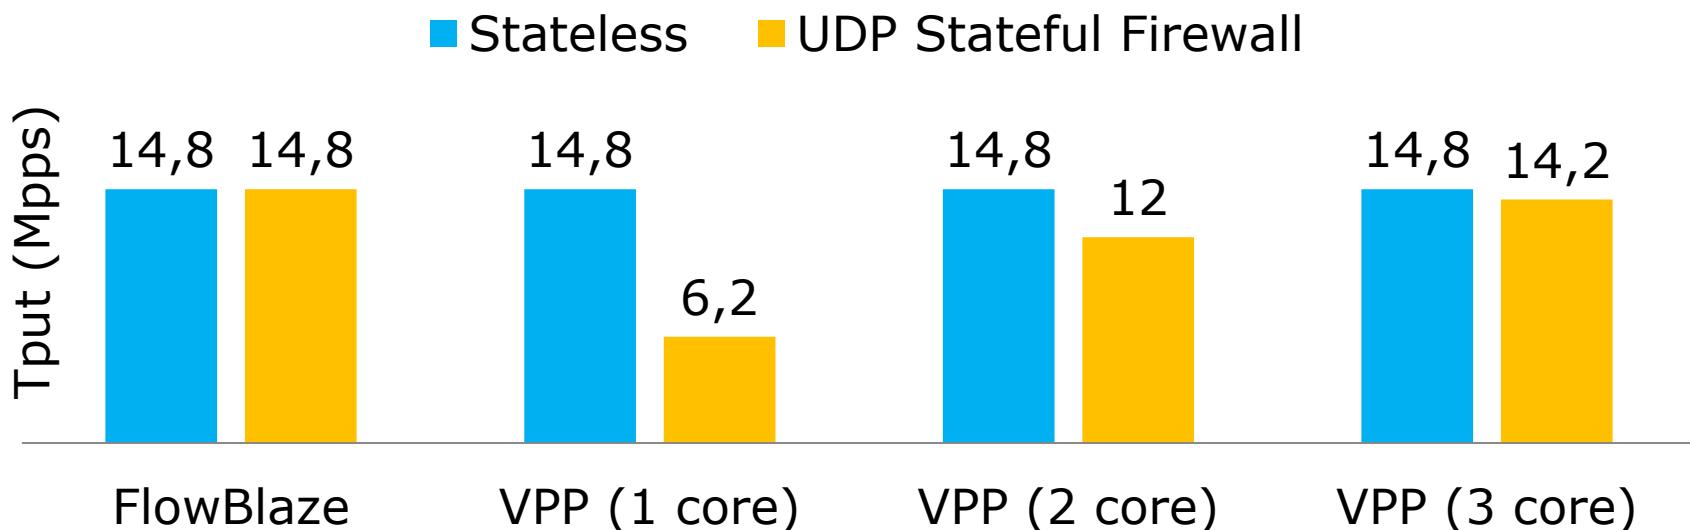
Use case
Server Load Balancer
UDP Stateful Firewall
Port Knocking Firewall
Flowlet load balancer
Traffic Policer
Big Flow Detector
SYN flood Detection and Mitigation
TCP optimistic ACK detection
TCP super spreader detection
Dynamic NAT
vEPC subscriber's quota verification
Switch Paxos Coordinator
Switch Paxos Acceptor
In-network KVS cache

FlowBlaze provides the same performance for all use cases

Test:
10Gb/s@64B
flow definition: 5-tuple

FlowBlaze:
NetFPGA@156.25MHz

Compared to:
DPDK-VPP on Xeon
X3470@2.93GHz, Intel
82599 10GbE NIC

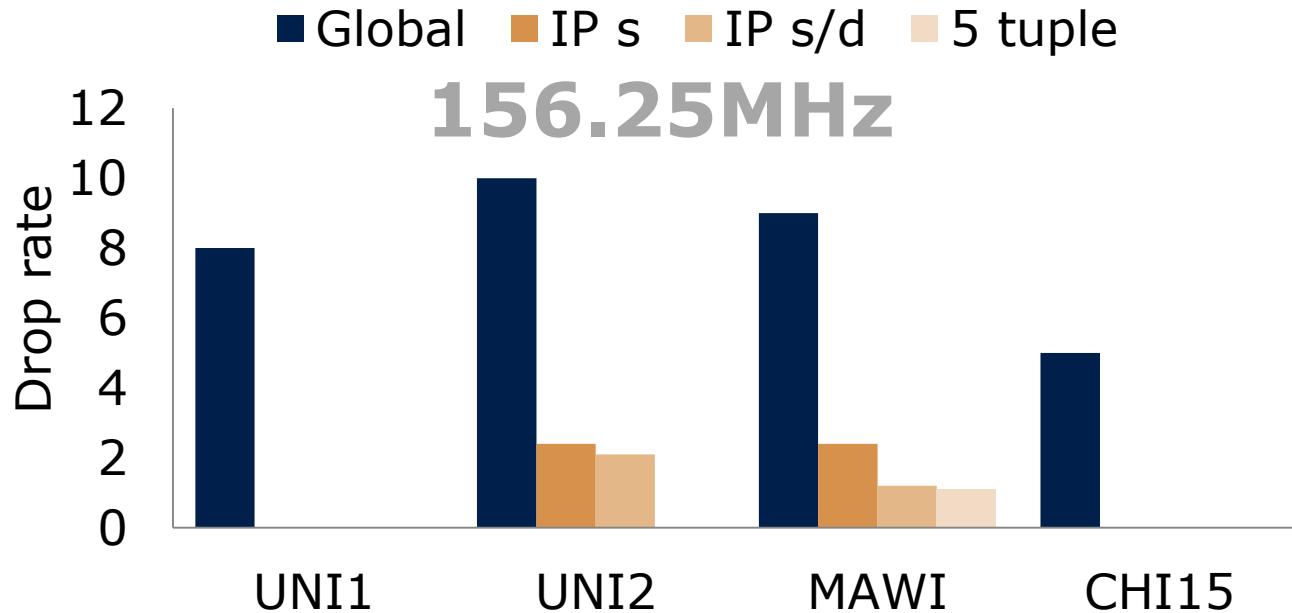


Stress test

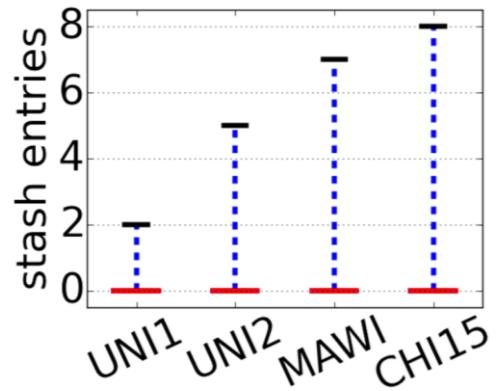
**Test:
40Gb/s@64B
(NetFPGA line rate)**

Trace	Max # active flows			Max # new flows/ms		
	IP s	IP s,d	5 tpl	IP s	IP s,d	5 tpl
UNI1	575	997	4k	13	19	39
UNI2	948	3k	7k	20	42	42
MW15	12k	130k	152k	38	112	114
CHI15	92k	147k	178k	135	144	144

Flow distributions



Stash



FlowBlaze

- **FSM Abstraction for packet processing**
- **Efficient FPGA implementation**

Benefits

- **Can keep state for 100Ks flows in flow tables**
- **Save several CPU cores for stateful NFs**
- **Power efficient (check the paper!)**
- **Low latency (check the paper!)**

Check the paper, there's a lot more!

FlowBlaze is open

Both software and hardware implementations
maintained by



Axbryd

<https://github.com/axbryd/FlowBlaze>

Thank you!
visit us and check our demo
at the poster session

\Orchestrating a brighter world

NEC