



# FlowBlaze.p4: a library for quick prototyping of stateful SDN applications in P4

*Daniele Moro\**, *Davide Sanvito<sup>^</sup>*, *Antonio Capone\**

\* Politecnico di Milano, Italy  
^ NEC Laboratories Europe, Germany

IEEE Conference on Network Function Virtualization and Software Defined Networks 2020  
9-12 November 2020  
Virtual Conference

# Outline


- Introduction
- The FlowBlaze architecture
- FlowBlaze.p4 library
- Using FlowBlaze.p4: Packet Limiter
- Use Case: Multi-Class Rate Limiter
- Conclusions

# Introduction - 1



5G and Mobile Edge Computing requires offloading of network functions to data plane

- **P4**: reference language for data plane programming
- **State Machine**: powerful abstraction to develop stateful packet processing
- **FlowBlaze**: EFSM-based stateful packet processing architecture



# Introduction - 2


## Issues with FlowBlaze utilization:

- Missing prototyping platform
- Manual (error-prone) mapping from EFSM to FlowBlaze table entries
- No FlowBlaze P4 implementation





## FlowBlaze.p4:

- FlowBlaze library implementation in P4
- Open source library
- GUI to automatically translate EFSM into table entries
- Exploit all the tools from the P4 Community



# The FlowBlaze\* architecture - 1

- EFSM based stateful packet processing
- Multi-stage: stateless (OpenFlow like) + stateful (EFSM-based) stages





\* S Pontarelli, et al. "Flowblaze: Stateful packet processing in hardware", USENIX NSDI 2019

# The FlowBlaze architecture - 2

## Stateful Stage

- Arbitrary Flow Definition with associated context (State and Flow Data Variable - FDV )
- Conditions evaluated on the FDVs ( $<$ ,  $>$ ,  $\leq$ ,  $\geq$ ,  $=$ ,  $\neq$ )
- Transitions in EFSM Table (if  $\langle$ conditions $\rangle$  and  $\langle$ state $\rangle$  then  $\langle$ new\_state, actions $\rangle$ )
- Update FDVs (+, -, \*,  $\langle\langle$ ,  $\rangle\rangle$ ) and packet action (e.g., forward, drop...)




# FlowBlaze.p4 library

## Stateful FlowBlaze Stage to P4

Communication between blocks via P4 Packet Metadata

Target: BMv2 software switch



Flow Context	<b>Registers*</b> *Indexing with hash function on the Flow Definition
Conditions	Series of if on the Flow Data Variables
EFSM Table	<b>Match Action Table</b> <b>Match:</b> <conditions results, state, arbitrary fields> <b>Action:</b> <Set new state, Set Packet Action ID, Set Update Functions>
Packet Action	<b>Match Action Table</b> Match: Packet Action ID Actions: User Defined P4 Actions
Update Functions	Series of if to map the set update function to the actual action on the FDV

# Using FlowBlaze.p4

Packet Limiter: EFSM

Flow Definition

- IPv4 source address

- State **0**

All the flows starts from this state.

Transition to the **Count** state.

- State **Count**

Counts the number of packets.


Auto-transition that count the packets when below the threshold

Transition to the **Drop** state when reached the threshold (10) packets.

- State **Drop**

Drop all the packets.

“Black-hole” state.



# Using FlowBlaze.p4

## Packet Limiter: Compile-time Configuration

### 1. Add FlowBlaze into your P4 application

### 2. #define:

- **Flow Scope:** source IP address
- **Packet actions:**
  - forward()
  - drop()
- **EFSM header match**
- **Condition header**

### 3. Compile the program

```
#include "../flowblaze_lib/flowblaze_metadata.p4"
#include "headers.p4"
#include "metadata.p4"
#include "../flowblaze_lib/flowblaze.p4"
...
apply {
    if (hdr.ethernet.isValid()) {
        FlowBlaze.apply(hdr, meta, standard_metadata);
        t_l2_fwd.apply();
    }
}

#define FLOW_SCOPE { hdr.ipv4.srcAddr }
#define CUSTOM_ACTIONS_DEFINITION @name(".FlowBlaze.forward") \
                                action forward() { \
                                    \
                                } \
                                @name(".FlowBlaze.drop") \
                                action drop() { \
                                    mark_to_drop(standard_metadata); \
                                    exit; \
                                }
#define CUSTOM_ACTIONS_DECLARATION forward; drop;
// Configuration parameter left black because not needed
//     #define METADATA_OPERATION_COND
//     #define EFSM_MATCH_FIELDS
//     #define CONTEXT_TABLE_SIZE
```

} Not needed for  
this example


# Using FlowBlaze.p4 (continued)

## Packet Limiter: Run-time Configuration

### 4. Run the GUI:

- Add states as in the EFSM
- Build the transitions as in the drawn EFSM:
  - Match + Condition
  - Update Function + Packet Action

### 5. Run in Mininet with the provided Docker infrastructure



# Use Cases

## Multi-Class Rate Limiter

### Flow Definition

- IPv4 source address

- State **0**

All the flows starts from this state.

“Classify” the traffic by the source IP address  
setting the *max\_bytes* FDV

- State **Allow**


Let traffic through, counting the bytes

Transition to the **Drop** state if in the predefined  
time-slot more than *max\_bytes* passed.

- State **Drop**

Drop all the packets.

Transition back to the **Allow** state when time-slot ends.




**Time-based transition:**  
@now is the packet ingress timestamp

# Conclusions

- FlowBlaze: abstraction for stateful packet processing based on EFSM
- Missing: prototyping platform and P4 implementation

We provide:

- **FlowBlaze.p4**: a library for prototyping with FlowBlaze and P4
- **GUI**: automatic translation of EFSM into runtime configuration
- **Docker-based environment**
- **Open source**: available on GitHub



Future works:

- ONOS Integration
- DC-style fabric integration (e.g., Trellis)



Daniele Moro  
daniele.moro@polimi.it



ADVANCED  
NETWORK  
TECHNOLOGIES  
LAB



<https://github.com/ANTLab-polimi/flowblaze.p4>

**Demo later today!**  
*“Demonstrating FlowBlaze.p4: fast prototyping for EFSM-based data plane applications”*



ADVANCED  
NETWORK  
TECHNOLOGIES  
LAB